

# Chapter ~~3~~ 6

## Deadlocks

But crises and deadlocks when they occur have at least this advantage, that they force us to think.

— *Jawaharlal Nehru*

# Overview

---

- ✦ Resources
- ✦ Why do deadlocks occur?
- ✦ Dealing with deadlocks
  - Ignoring them: ostrich algorithm
  - Detecting & recovering from deadlock
  - Avoiding deadlock
  - Preventing deadlock

# Resources

---

- ♦ Resource: something a process uses
  - Usually limited (at least somewhat)
- ♦ Examples of computer resources
  - Printers
  - Semaphores / locks
  - Memory
  - Tables (in a database)
- ♦ Processes need access to resources in reasonable order
- ♦ Two types of resources:
  - Preemptable resources: can be taken away from a process with no ill effects
  - Nonpreemptable resources: will cause the process to fail if taken away


# Using resources

---

- ✦ Sequence of events required to use a resource
  - Request the resource
  - Use the resource
  - Release the resource
- ✦ Can't use the resource if request is denied
  - Requesting process has options
    - Block and wait for resource
    - Continue (if possible) without it: may be able to use an alternate resource
    - Process fails with error code
  - Some of these may be able to prevent deadlock...

# When do deadlocks happen?

- ✦ Suppose
  - Process 1 holds resource A and requests resource B
  - Process 2 holds B and requests A
  - Both can be blocked, with neither able to proceed
- ✦ Deadlocks occur when ...
  - Processes are granted exclusive access to devices or software constructs (resources)
  - Each deadlocked process needs a resource held by another deadlocked process



# What is a deadlock?

---

- ◆ Formal definition:

“A set of processes is deadlocked if each process in the set is waiting for an event that only another process in the set can cause.”



- ◆ Usually, the event is release of a currently held resource
- ◆ In deadlock, none of the processes can
  - Run
  - Release resources
  - Be awakened

# Four conditions for deadlock

---

- ✦ Mutual exclusion
  - Each resource is assigned to at most one process
- ✦ Hold and wait
  - A process holding resources can request more resources
- ✦ No preemption
  - Previously granted resources cannot be forcibly taken away
- ✦ Circular wait
  - There must be a circular chain of 2 or more processes where each is waiting for a resource held by the next member of the chain

# Resource allocation graphs



- ✦ Resource allocation modeled by directed graphs
- ✦ Example 1:
  - Resource R assigned to process A
- ✦ Example 2:
  - Process B is requesting / waiting for resource S
- ✦ Example 3:
  - Process C holds T, waiting for U
  - Process D holds U, waiting for T
  - C and D are in deadlock!



# Dealing with deadlock



---

- ✦ How can the OS deal with deadlock?
  - Ignore the problem altogether!
    - Hopefully, it'll never happen...
  - Detect deadlock & recover from it
  - Dynamically avoid deadlock
    - Careful resource allocation
  - Prevent deadlock
    - Remove at least one of the four necessary conditions
- ✦ We'll explore these tradeoffs

# Getting into deadlock



A

Acquire R  
Acquire S  
Release R  
Release S





B

Acquire S  
Acquire T  
Release S  
Release T



C

Acquire T  
Acquire R  
Release T  
Release R



# The Ostrich Algorithm

---

- ✦ Pretend there's no problem
- ✦ Reasonable if
  - Deadlocks occur very rarely
  - Cost of prevention is high
- ✦ UNIX and Windows take this approach
  - Resources (memory, CPU, disk space) are plentiful
  - Deadlocks over such resources rarely occur
  - Deadlocks typically handled by rebooting
- ✦ Trade off between convenience and correctness


# Not getting into deadlock...

- ✦ Many situations may result in deadlock (but don't have to)
  - In previous example, A could release R before C requests R, resulting in no deadlock
  - Can we always get out of it this way?
- ✦ Find ways to:
  - Detect deadlock and reverse it
  - Stop it from happening in the first place

# Detecting deadlocks using graphs

- ✦ Process holdings and requests in the table and in the graph (they're equivalent)
- ✦ Graph contains a cycle  $\Rightarrow$  deadlock!
  - Easy to pick out by looking at it (in this case)
  - Need to mechanically detect deadlock
- ✦ Not all processes are deadlocked (A, C, F not in deadlock)

Process	Holds	Wants
A	R	S
B		T
C		S
D	U	S,T
E	T	V
F	W	S
G	V	U



# Deadlock detection algorithm

- ✦ General idea: try to find cycles in the resource allocation graph
- ✦ Algorithm: depth-first search at each node
  - Mark arcs as they're traversed
  - Build list of visited nodes
  - If node to be added is already on the list, a cycle exists!
- ✦ Cycle  $\Rightarrow$  deadlock

```
For each node N in the graph {  
    Set L = empty list  
    unmark all arcs  
    Traverse (N,L)  
}  
If no deadlock reported by now,  
there isn't any  
  
define Traverse (C,L) {  
    If C in L, report deadlock!  
    Add C to L  
    For each unmarked arc from C {  
        Mark the arc  
        Set A = arc destination  
        /* NOTE: L is a  
           local variable */  
        Traverse (A,L)  
    }  
}
```

# Recovering from deadlock

---

- ✦ Recovery through preemption
  - Take a resource from some other process
  - Depends on nature of the resource and the process
- ✦ Recovery through rollback
  - Checkpoint a process periodically
  - Use saved state to restart the process if it's in deadlock
  - May present a problem if the process affects lots of “external” things
- ✦ Recovery through killing processes
  - Crudest but simplest way to break a deadlock: kill one of the processes in the deadlock cycle
  - Other processes can get its resources
  - Try to choose a process that can be rerun from the start
    - Pick one that hasn't run too far already

# Preventing deadlock

---

- ✦ Deadlock can be completely prevented!
- ✦ Ensure that at least one of the conditions for deadlock never occurs
  - Mutual exclusion
  - Circular wait
  - Hold & wait
  - No preemption
- ✦ Not always possible...



# Eliminating mutual exclusion

---

- ✦ Some devices (such as printer) can be spooled
  - Only the printer daemon uses printer resource
  - This eliminates deadlock for printer
- ✦ Not all devices can be spooled
- ✦ Principle:
  - Avoid assigning resource when not absolutely necessary
  - As few processes as possible actually claim the resource


# Attacking “hold and wait”

---

- ✦ Require processes to request resources before starting
  - A process never has to wait for what it needs
- ✦ This can present problems
  - A process may not know required resources at start of run
  - This also ties up resources other processes could be using
    - Processes will tend to be conservative and request resources they might need
- ✦ Variation: a process must give up all resources before making a new request
  - Process is then granted all prior resources as well as the new ones
  - Problem: what if someone grabs the resources in the meantime—how can the process save its state?


# Attacking “no preemption”

- ✦ This is not usually a viable option
- ✦ Consider a process given the printer
  - Halfway through its job, take away the printer
  - Confusion ensues!
- ✦ May work for some resources
  - Forcibly take away memory pages, suspending the process
  - Process may be able to resume with no ill effects



# Attacking “circular wait”

- ✦ Assign an order to resources
- ✦ Always acquire resources in numerical order
  - Need not acquire them all at once!
- ✦ Circular wait is prevented
  - A process holding resource  $n$  can't wait for resource  $m$  if  $m < n$
  - No way to complete a cycle
    - Place processes above the highest resource they hold and below any they're requesting
    - All arrows point up!



# Deadlock prevention: summary

Condition	Prevented by
Mutual exclusion	Spool everything
Hold and wait	Request all resources initially
No preemption	Take resources away if there's not a complete set
Circular wait	Order resources numerically

# Example: two-phase locking

---

## ◆ Phase One

- Process tries to lock all data it needs, one at a time
- If needed data found locked, start over  
(no real work done in phase one)

## ◆ Phase Two

- Perform updates
- Release locks

## ◆ Note similarity to requesting all resources at once

## ◆ This is often used in databases

## ◆ What condition does this avoid?

# “Non-resource” deadlocks

---

- ✦ Possible for two processes to deadlock
  - Each is waiting for the other to do some task
- ✦ Can happen with semaphores
  - Each process required to do a down() on two semaphores (mutex and another)
  - If done in wrong order, deadlock results
- ✦ Semaphores could be thought of as resources...

# Handling resources with multiple instances

---

- ✦ Previous algorithm only works if there's one instance of each resource
- ✦ If there are multiple instances of each resource, we need a different method
  - Track current usage and requests for each process
  - To detect deadlock, try to find a scenario where all processes can finish
  - If no such scenario exists, we have deadlock



# Deadlock detection algorithm

	A	B	C	D
Avail	2	3	0	1

Hold

Process	A	B	C	D
1	0	3	0	0
2	1	0	1	1
3	0	2	1	0
4	2	2	3	0

Want

Process	A	B	C	D
1	3	2	1	0
2	2	2	0	0
3	3	5	3	1
4	0	4	1	1

```

current = avail;
for (j = 0; j < N; j++) {
    for (k=0; k<N; k++) {
        if (finished[k])
            continue;
        if (want[k] < current) {
            finished[k] = 1;
            current += hold[k];
            break;
        }
    }
    if (k==N) {
        printf "Deadlock!\n";
        // finished[k]==0 means
        // process
        // is in the deadlock
        break;
    }
}
    
```

Note: want[j],hold[j],current,avail are arrays!

# Safe and unsafe states

Has Max			Has Max			Has Max			Has Max			Has Max		
A	3	9	A	3	9	A	3	9	A	3	9	A	3	9
B	2	4	B	4	4	B	0	-	B	0	-	B	0	-
C	2	7	C	2	7	C	2	7	C	7	7	C	0	-
Free: 3			Free: 1			Free: 5			Free: 0			Free: 7		

Demonstration that the first state is safe

Has Max			Has Max			Has Max			Has Max		
A	3	9	A	4	9	A	4	9	A	4	9
B	2	4	B	2	4	B	4	4	B	0	-
C	2	7	C	2	7	C	2	7	C	2	7
Free: 3			Free: 2			Free: 0			Free: 4		

Demonstration that the second state is unsafe

# Banker's Algorithm for a single resource

	Has	Max
A	0	6
B	0	5
C	0	4
D	0	7

Free: 10

Any sequence finishes

	Has	Max
A	1	6
B	1	5
C	2	4
D	4	7

Free: 2

C,B,A,D finishes

	Has	Max
A	1	6
B	2	5
C	2	4
D	4	7

Free: 1

Deadlock (unsafe state)

- ✦ Bankers' algorithm: before granting a request, ensure that a sequence exists that will allow all processes to complete
  - Use previous methods to find such a sequence
  - If a sequence exists, allow the requests
  - If there's no such sequence, deny the request
- ✦ Can be slow: must be done on each request!

# Banker's Algorithm with multiple resources

	Process	Tape drives	Plotters	Scanners	DVD-ROMs
A	3	0	1	1	
B	0	1	0	0	
C	1	1	1	0	
D	1	1	0	1	
E	0	0	0	0	

Resources assigned

	Process	Tape drives	Plotters	Scanners	DVD-ROMs
A	1	1	0	0	
B	0	1	1	2	
C	3	1	0	0	
D	0	0	1	0	
E	2	1	1	0	

Resources still needed

<b>Total resources:</b>	6	3	4	2
<b>Already claimed:</b>	5	3	2	2
<b>Available:</b>	1	0	2	0

# Starvation

---

- ✦ Algorithm to allocate a resource
  - Give the resource to the shortest job first
- ✦ Works great for multiple short jobs in a system
- ✦ May cause long jobs to be postponed indefinitely
  - Even though not blocked
- ✦ Solution
  - First-come, first-serve policy
- ✦ Starvation can lead to deadlock
  - Process starved for resources can be holding resources
  - If those resources aren't used and released in a timely fashion, shortage could lead to deadlock