

# Processes, Threads, and Scheduling



# Processes and threads

---

- ◆ Processes
- ◆ Threads
- ◆ Scheduling
- ◆ Interprocess communication (IPC)
- ◆ Classical IPC problems

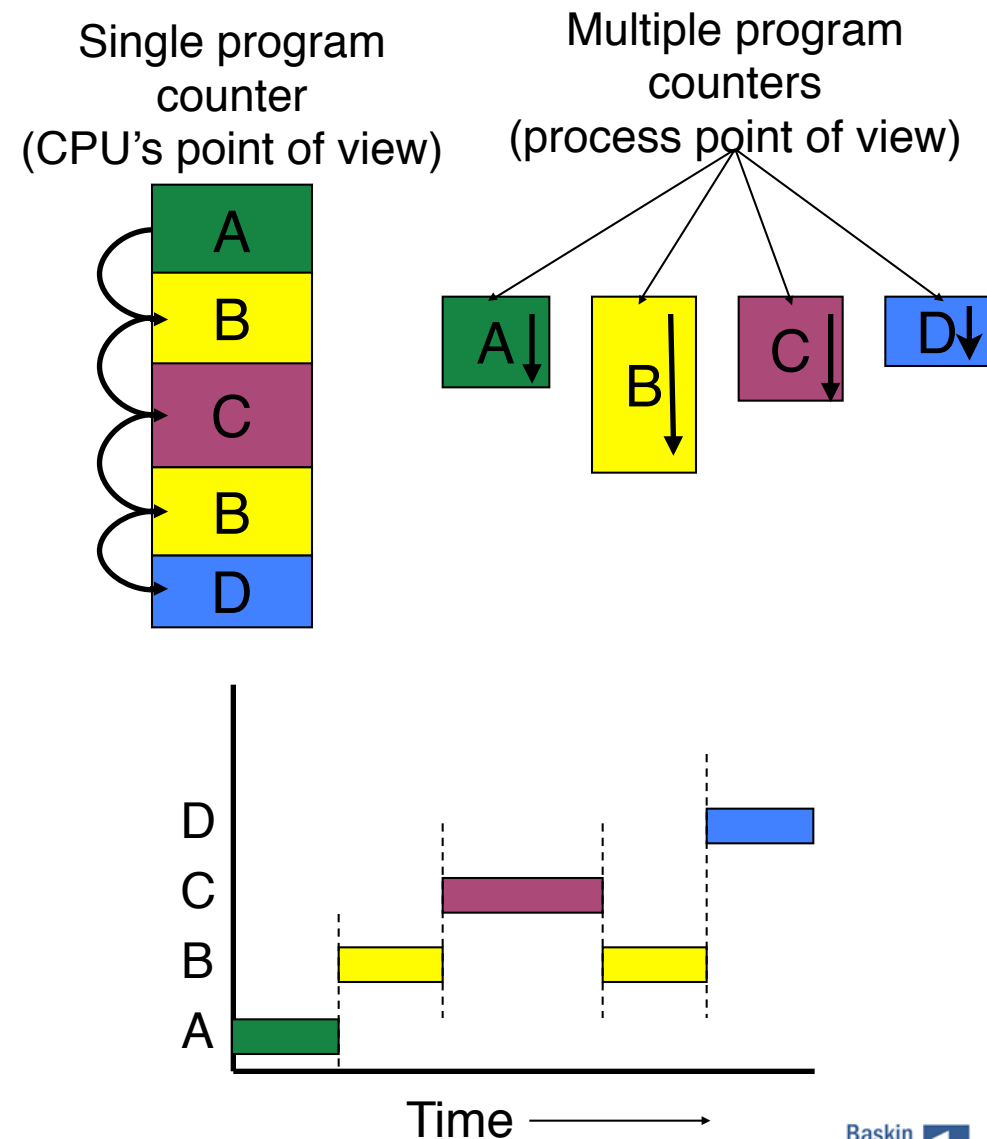
# What is a process?

---

- ✦ Code, data, and stack
  - Usually (but not always) has its own address space
- ✦ Program state
  - CPU registers
  - Program counter (current location in the code)
  - Stack pointer
- ✦ Only one process can be running in the CPU at any given time!

# The process model

- ◆ Multiprogramming of four programs
- ◆ Conceptual model
  - 4 independent processes
  - Processes run sequentially
- ◆ Only one program active at any instant!
  - That instant can be very short...
  - Only applies if there's a single CPU (with a single core) in the system



# When is a process created?

---

- ✦ Processes can be created in two ways
  - System initialization: one or more processes created when the OS starts up
  - Execution of a process creation system call: something explicitly asks for a new process
- ✦ System calls can come from
  - User request to create a new process (system call executed from user shell)
  - Already running processes
    - User programs
    - System daemons

# When do processes end?

---

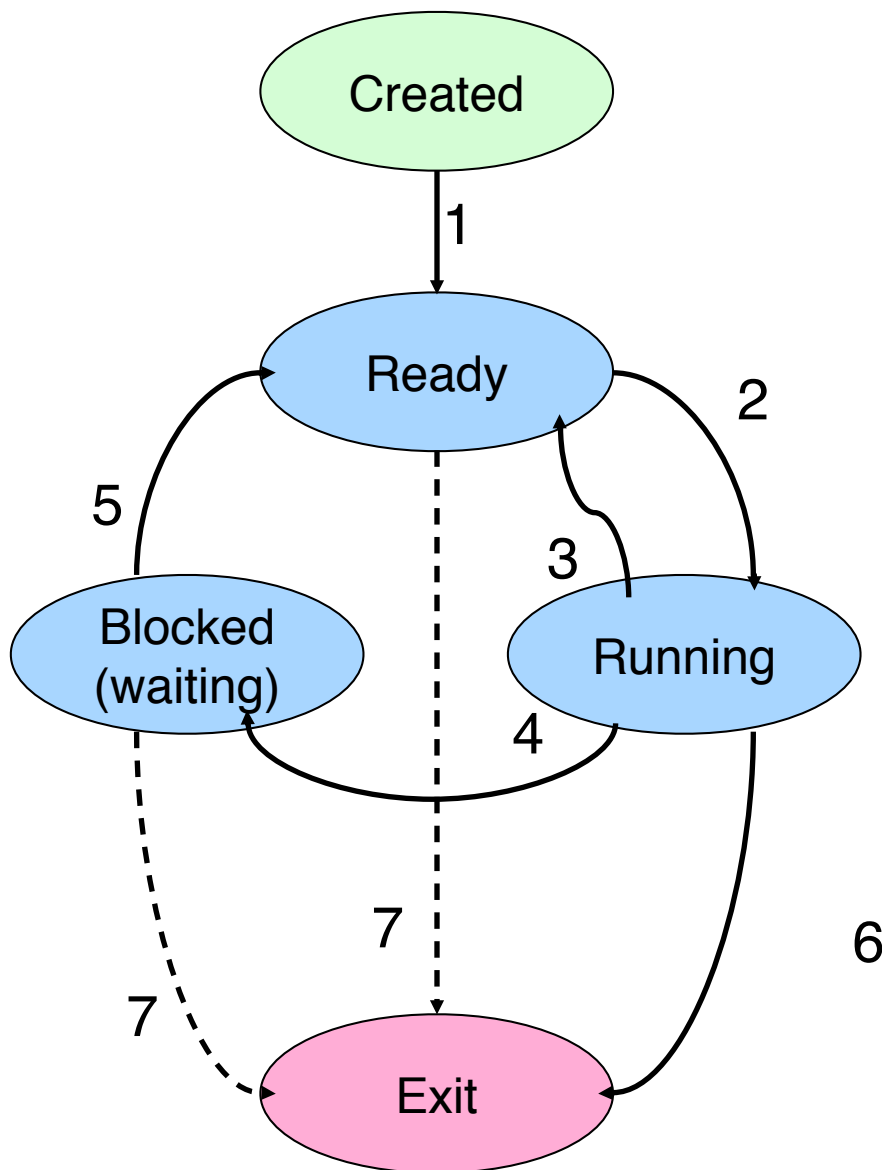
- ✦ Conditions that terminate processes can be
  - Voluntary
  - Involuntary
- ✦ Voluntary
  - Normal exit
  - Error exit
- ✦ Involuntary
  - Fatal error (only sort of involuntary)
  - Killed by another process

# Process hierarchies

---

- ✦ Parent creates a child process
  - Child processes can create their own children
- ✦ Forms a hierarchy
  - UNIX calls this a “process group”
  - If a process terminates, its children are “inherited” by the terminating process’s parent
- ✦ Windows has process groups
  - Multiple processes grouped together
  - One process is the “group leader”

# Process states



## ◆ Process in one of 5 states

- Created
- Ready
- Running
- Blocked
- Exit

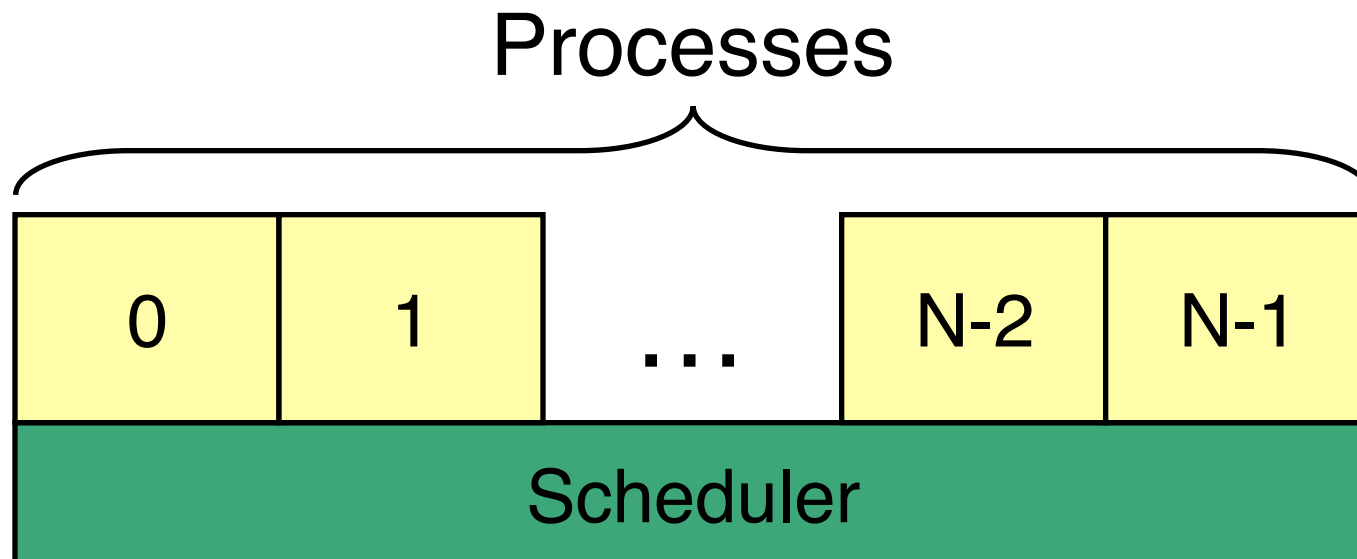
## ◆ Transitions between states

- Process enters ready queue
- Scheduler picks this process
- Scheduler picks a different process
- Process waits for event (such as I/O)
- Event occurs
- Process exits
- Process ended by another process

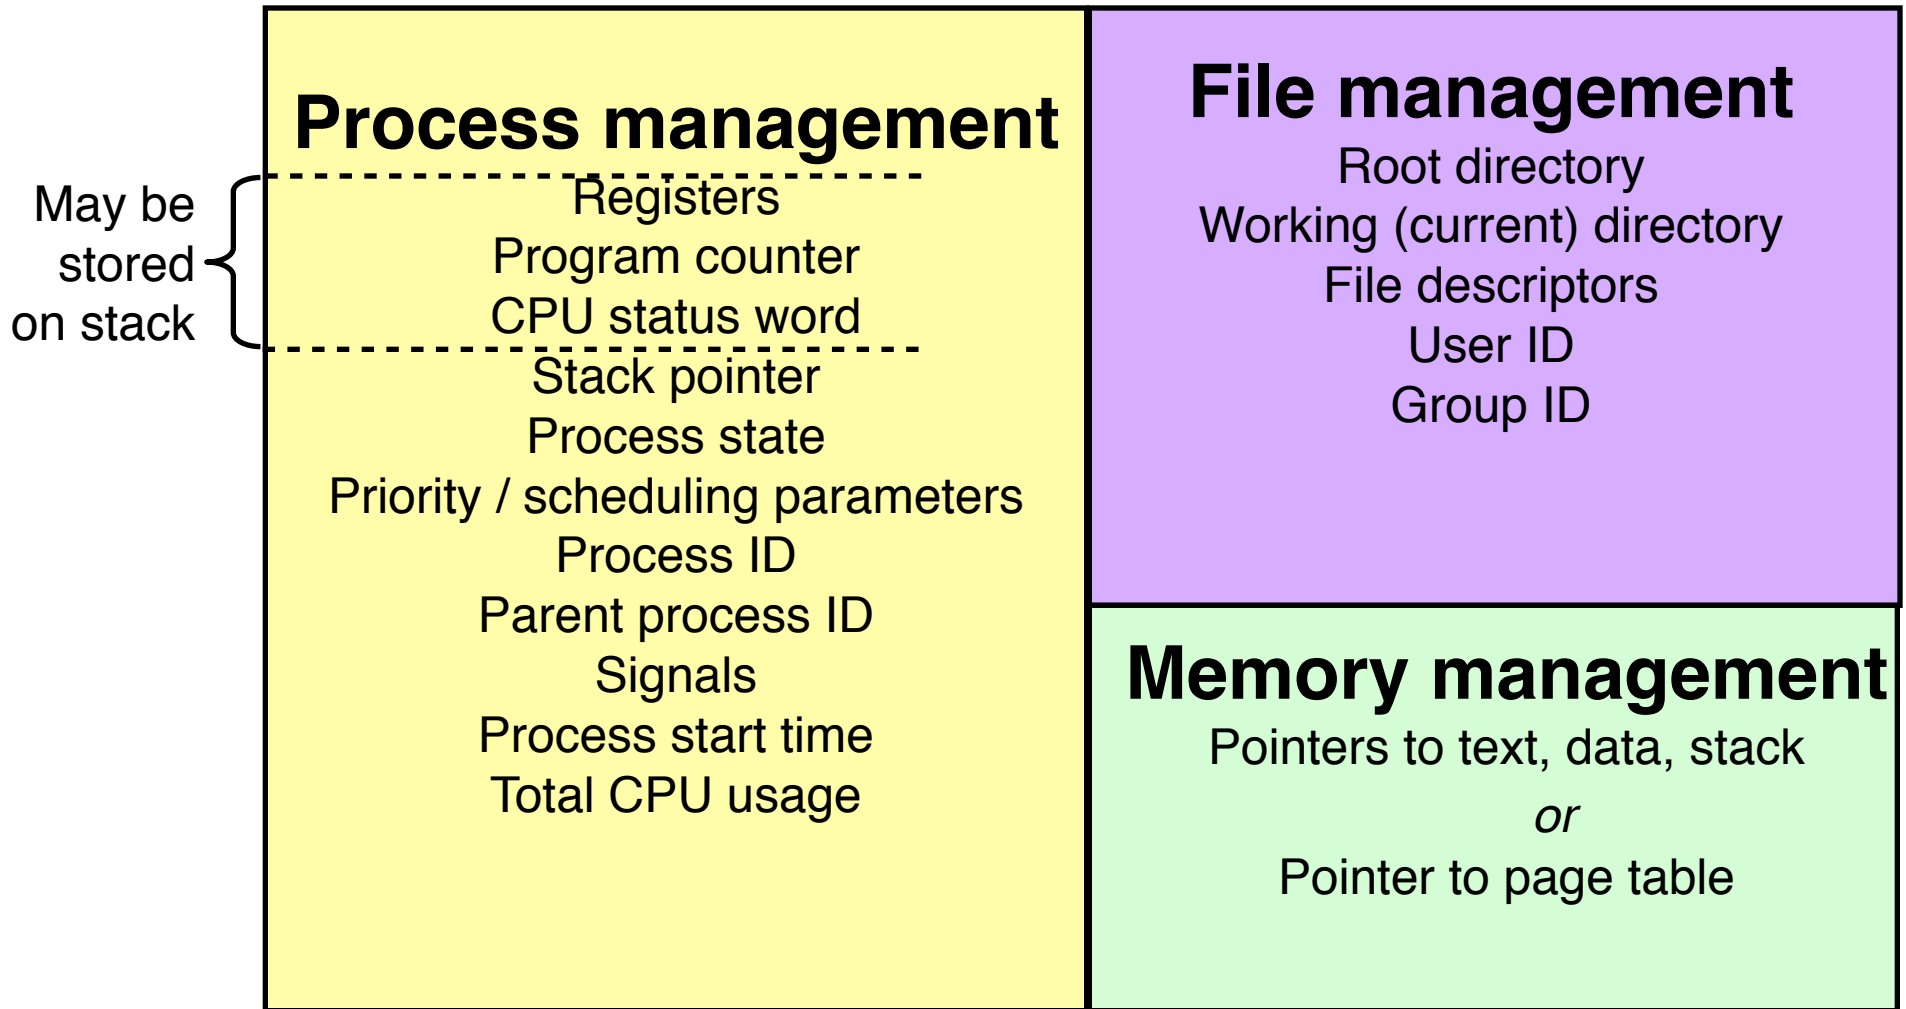


# Processes in the OS

- ◆ Two “layers” for processes
- ◆ Lowest layer of process-structured OS handles interrupts, scheduling
- ◆ Above that layer are sequential processes
  - Processes tracked in the process table
  - Each process has a process table entry



# What's in a process table entry?




# What happens on a trap/ interrupt?

---

1. Hardware saves program counter (on stack or in a special register)
2. Hardware loads new PC, identifies interrupt
3. Assembly language routine saves registers
4. Assembly language routine sets up stack
5. Assembly language calls C to run service routine
6. Service routine calls scheduler
7. Scheduler selects a process to run next (might be the one interrupted...)
8. Assembly language routine loads PC & registers for the selected process

# Threads: “processes” sharing memory

- ◆ Process == address space
- ◆ Thread == program counter / stream of instructions
- ◆ Two examples
  - Three processes, each with one thread
  - One process with three threads



# Process & thread information

## Per process items

Address space  
Open files  
Child processes  
Signals & handlers  
Accounting info  
*Global variables*

## Per thread items

Program counter  
Registers  
Stack & stack pointer  
State


## Per thread items

Program counter  
Registers  
Stack & stack pointer  
State

## Per thread items

Program counter  
Registers  
Stack & stack pointer  
State

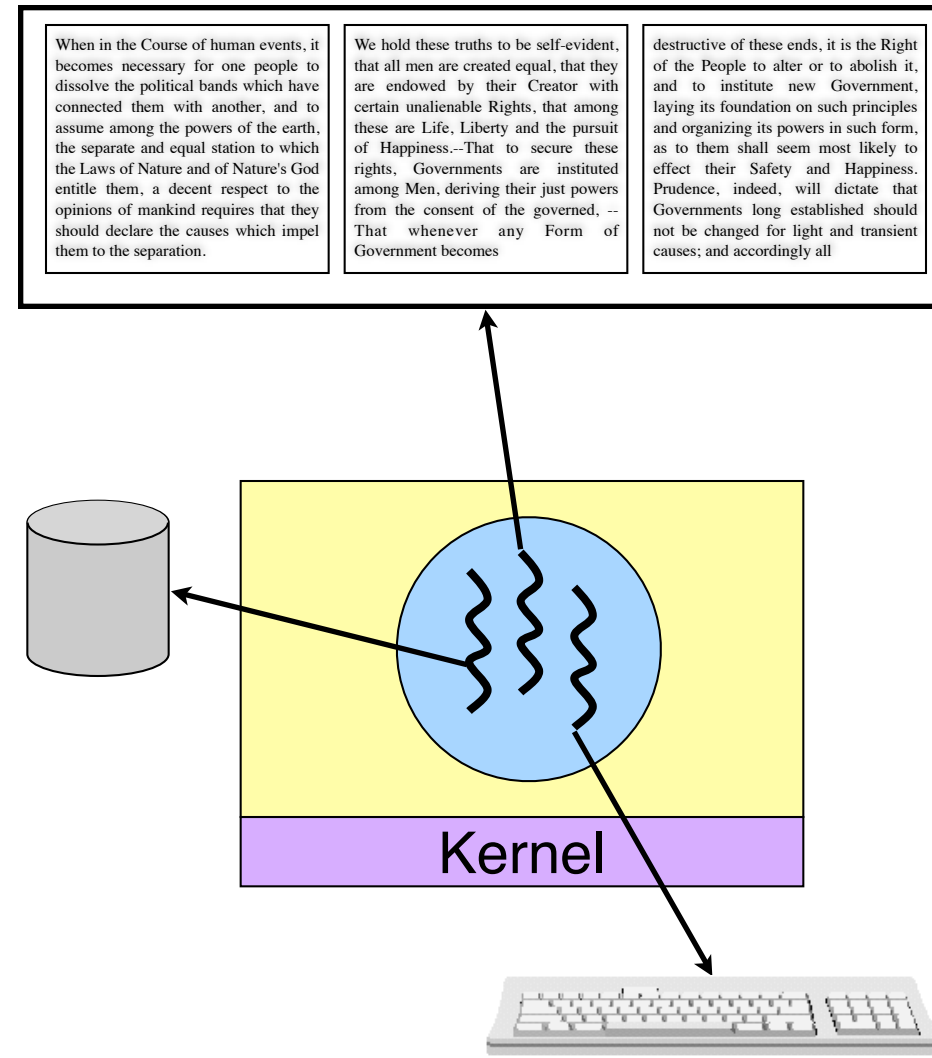
# Threads & stacks



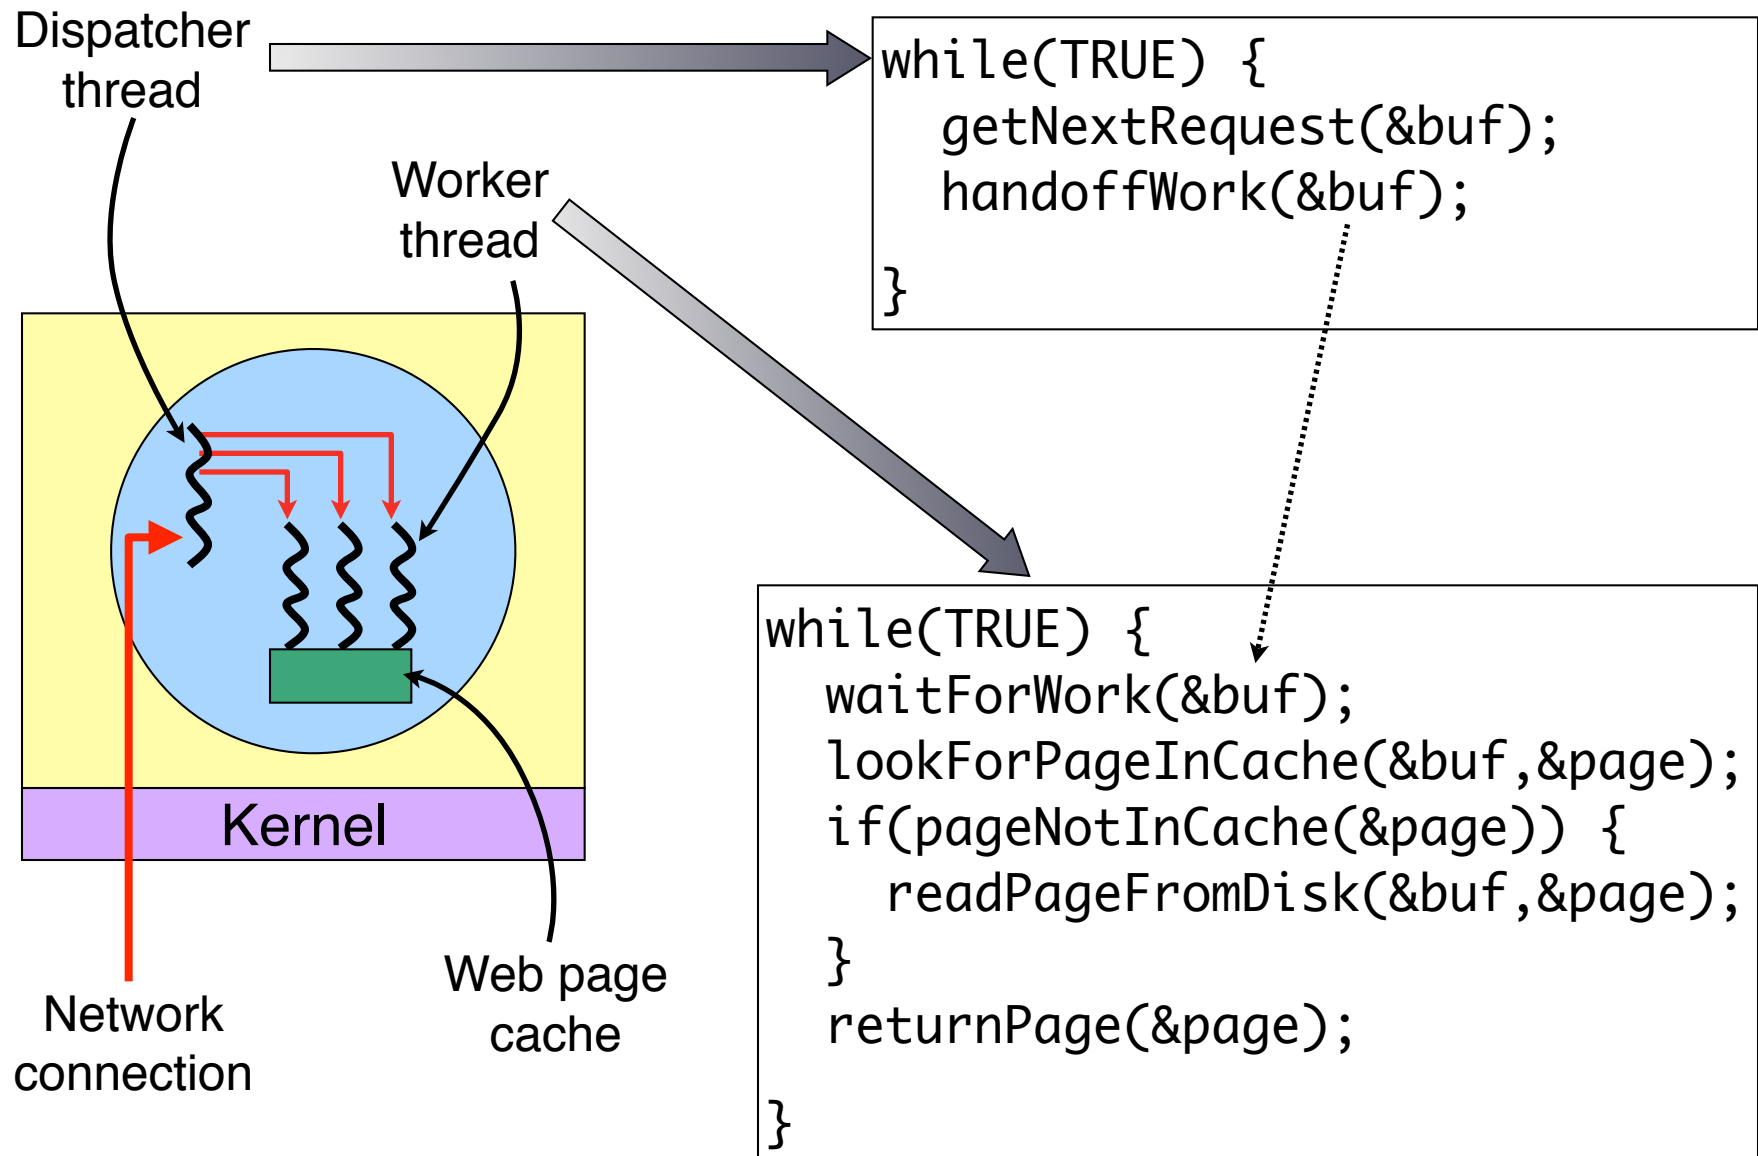
➡ Each thread has its own stack!

# Why use threads?

- ◆ Allow a single application to do many things at once
  - Simpler programming model
  - Less waiting
- ◆ Threads are faster to create or destroy
  - No separate address space
- ◆ Overlap computation and I/O
  - Could be done without threads, but it's harder
- ◆ Example: word processor
  - Thread to read from keyboard
  - Thread to format document
  - Thread to write to disk



# Multithreaded Web server





# Three ways to build a server

---


- ✦ Multithreaded server
  - Parallelism
  - Blocking system calls
  - May use pop-up threads: create a new thread in response to an incoming message (rather than reusing a thread)
- ✦ Single-threaded process: slow, but easier to do
  - No parallelism
  - Blocking system calls
- ✦ Finite-state machine (event model)
  - Each activity has its own state: states change when system calls complete or interrupts occur
  - Parallelism
  - Nonblocking system calls

# Issues with using threads

---

- ✦ May be tricky to convert single-threaded code to multithreaded code
- ✦ Re-entrant code
  - Code must function properly when multiple threads are using it simultaneously
  - Need to be careful when using static or global variables
    - Returned structures
    - Buffers
- ✦ Error management
  - What happens when just a single thread has an error?
  - Can't simply kill the process, since other threads might be running

# Implementing threads



## User-level threads

- + No need for kernel support
- May be slower than kernel threads
- Harder to do non-blocking I/O

## Kernel-level threads

- + More flexible scheduling
- + Non-blocking I/O
- Not portable

# POSIX threads

- ✦ Standard interface to threading library
- ✦ May be implemented in either user or kernel space
  - Some operating systems provide support for both!
- ✦ Allows thread-based programs to be portable

Thread call (Pthread_xx)	Description
create	Create a new thread
exit	Terminate the calling thread
join	Wait for a specific thread to exit
yield	Release the CPU, allowing another thread to run



# Processes & threads in Linux

---

- ✦ Linux supports kernel-level threads (lightweight processes)
  - Share address space, file descriptors, etc.
  - Each has its own process descriptor in memory
- ✦ Linux processes (incl. lightweight) all have unique identifiers
  - Threads sharing address space are grouped into *process groups*
  - Identifier shared by the group is that of the *leader*

# Linux process information

---

- ✦ Each process has its own 8KB region that stores
  - Kernel stack
    - Kernel has a small stack: about 4KB!
  - Low-level thread information
- ✦ Other information stored in a separate data structure
  - Memory allocated to the process
  - Open files
  - Signal information


# Scheduling

---

- ✦ What is scheduling?
  - Goals
  - Mechanisms
- ✦ Scheduling on batch systems
- ✦ Scheduling on interactive systems
- ✦ Other kinds of scheduling
  - Real-time scheduling

# Why schedule processes?

- ✦ Bursts of CPU usage alternate with periods of I/O wait
- ✦ Some processes are CPU-bound: they don't make many I/O requests
- ✦ Other processes are I/O-bound and make many kernel requests





# When are processes scheduled?

---

- ✦ At the time they enter the system
  - Common in batch systems
  - Two types of batch scheduling
    - Submission of a new job causes the scheduler to run
    - Scheduling only done when a job voluntarily gives up the CPU (i.e., while waiting for an I/O request)
- ✦ At relatively fixed intervals (clock interrupts)
  - Necessary for interactive systems
  - May also be used for batch systems
  - Scheduling algorithms at each interrupt, and picks the next process from the pool of “ready” processes

# Scheduling goals

---

## ✦ All systems

- Fairness: give each process a fair share of the CPU
- Enforcement: ensure that the stated policy is carried out
- Balance: keep all parts of the system busy

## ✦ Batch systems

- Throughput: maximize jobs per unit time (hour)
- Turnaround time: minimize time users wait for jobs
- CPU utilization: keep the CPU as busy as possible

## ✦ Interactive systems

- Response time: respond quickly to users' requests
- Proportionality: meet users' expectations

## ✦ Real-time systems

- Meet deadlines: missing deadlines is a system failure!
- Predictability: same type of behavior for each time slice

# Measuring scheduling performance

---

- ✦ Throughput
  - Amount of work completed per second (minute, hour)
  - Higher throughput usually means better utilized system
- ✦ Response time
  - Response time is time from when a command is submitted until results are returned
  - Can measure average, variance, minimum, maximum, ...
  - May be more useful to measure time spent waiting
  - Can also measure how often response time is faster than a given time (e.g., 100 ms): useful for real-time systems and servers
- ✦ Turnaround time
  - Like response time, but for batch jobs (response is the completion of the process)
- ✦ Usually not possible to optimize for *all* metrics with a single scheduling algorithm

# Interactive vs. batch scheduling

---

## **Batch**

First-Come-First-Served  
(FCFS)

Shortest Job First (SJF)

Shortest Remaining Time  
First (SRTF)

Priority (non-preemptive)

## **Interactive**

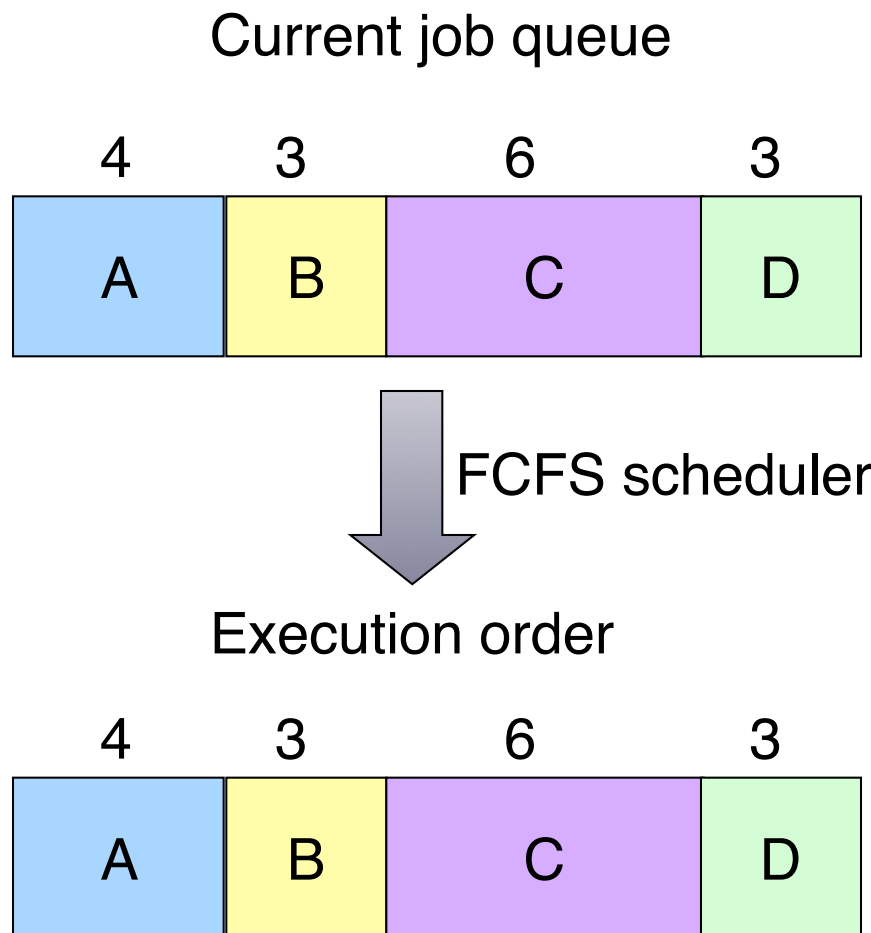
Round-Robin (RR)

Priority (preemptive)

Multi-level feedback queue


Lottery scheduling

# First Come, First Served (FCFS)



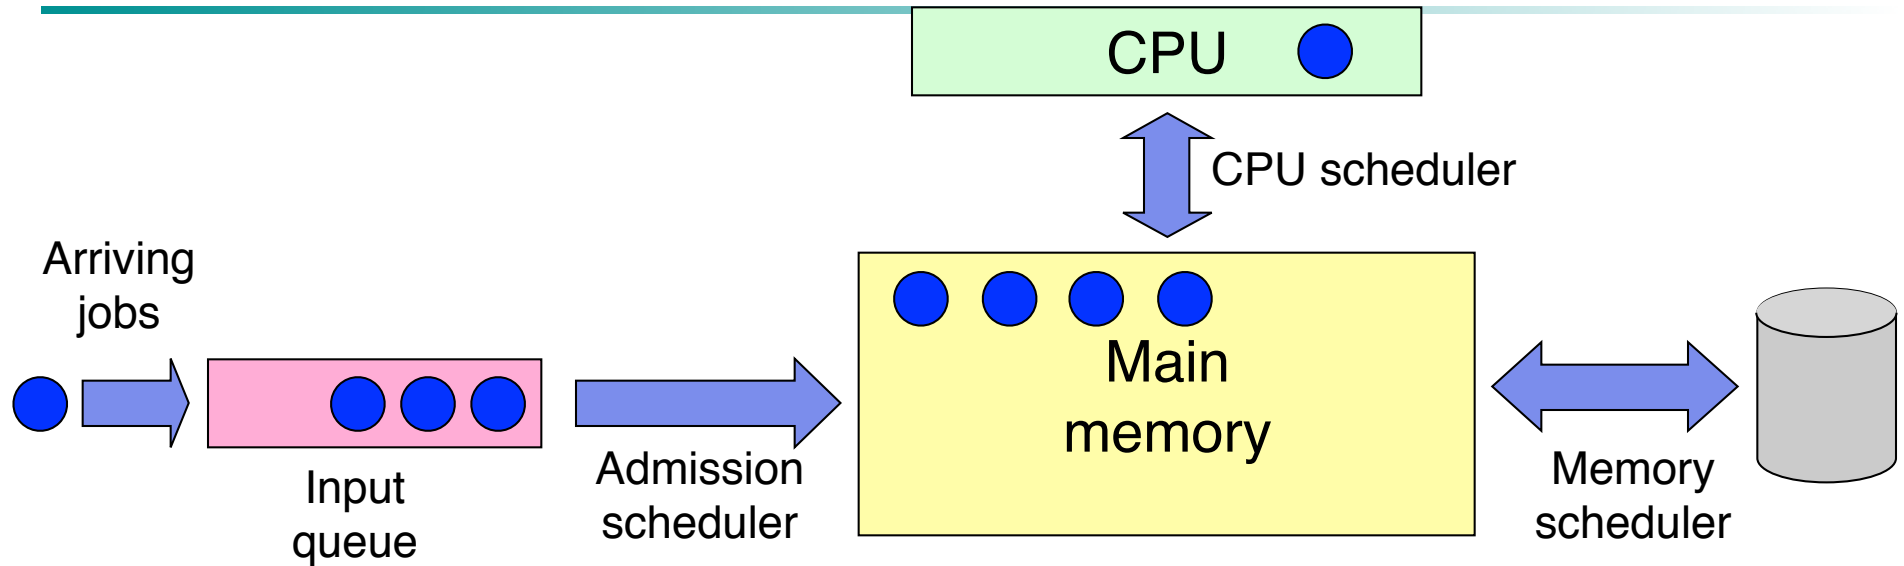
- ✦ Goal: do jobs in the order they arrive
  - Fair in the same way a bank teller line is fair
- ✦ Simple algorithm!
- ✦ Problem: long jobs delay every job after them
  - Many processes may wait for a single long job

# Shortest Job First (SJF)



- ✦ Goal: do the shortest job first
  - Short jobs complete first
  - Long jobs delay every job after them
- ✦ Jobs sorted in increasing order of execution time
  - Ordering of ties doesn't matter
- ✦ Shortest Remaining Time First (SRTF): preemptive form of SJF
  - Re-evaluate when a new job is submitted
- ✦ Problem: how does the scheduler know how long a job will take?

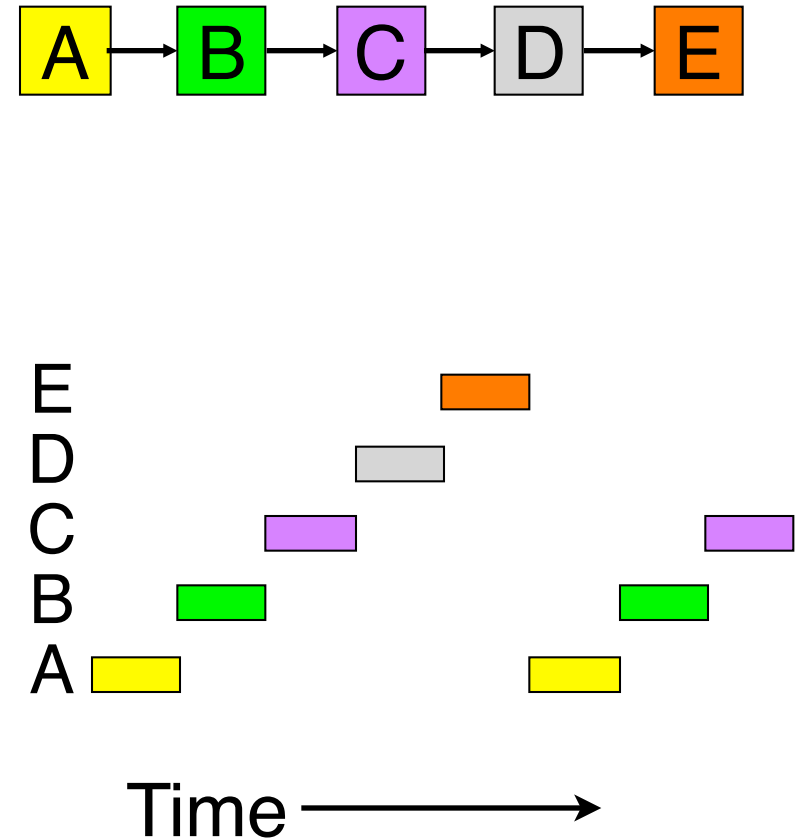
# Three-level scheduling



- ✦ Jobs held in input queue until moved into memory
  - Pick “complementary jobs”: small & large, CPU- & I/O-intensive
  - Jobs move into memory when admitted
- ✦ CPU scheduler picks next job to run
- ✦ Memory scheduler picks some jobs from main memory and moves them to disk if insufficient memory space

# Round Robin (RR) scheduling

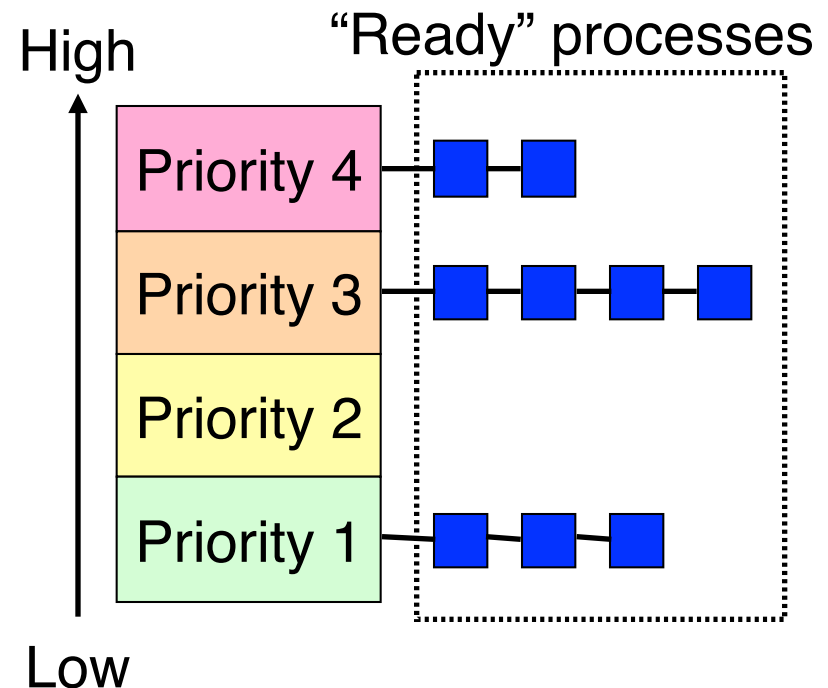
- ◆ Scheduling interactive processes
  - Give each process a fixed time slot (quantum)
  - Rotate through “ready” processes
  - Each process makes some progress
- ◆ What’s a good quantum?
  - Too short: many process switches hurt efficiency
  - Too long: poor response to interactive requests
  - Typical length: 10–100 ms
- ◆ “Strict” rotation: round robin





# Priority scheduling

- ◆ Assign a priority to each process
  - “Ready” process with highest priority allowed to run
  - Running process may be interrupted after its quantum expires
- ◆ Priorities may be assigned dynamically
  - Reduced when a process uses CPU time
  - Increased when a process waits for I/O
- ◆ Often, processes grouped into multiple queues based on priority, and run round-robin per queue



# Shortest process next

---

- ✦ Run the process that will finish the soonest
  - In interactive systems, job completion time is unknown!
- ✦ Guess at completion time based on previous runs
  - Update estimate each time the job is run
  - Estimate is a combination of previous estimate and most recent run time
- ✦ Not often used because round robin with priority works so well!

# Lottery scheduling


---

- ✦ Give processes “tickets” for CPU time
  - More tickets  $\Rightarrow$  higher share of CPU
- ✦ Each quantum, pick a ticket at random
  - If there are  $n$  tickets, pick a number from 1 to  $n$
  - Process holding the ticket gets to run for a quantum
- ✦ Over the long run, each process gets the CPU  $m/n$  of the time if the process has  $m$  of the  $n$  existing tickets
- ✦ Tickets can be transferred
  - Cooperating processes can exchange tickets
  - Clients can transfer tickets to a server so it can have a higher priority

# Scheduling in FreeBSD

- ✦ Quantum is 100 ms: longest that's OK for interactive scheduling
- ✦ Scheduler is based on multi-level feedback queues
  - Priority is based on two things
    - Resource requirements: blocked threads have higher priority when rescheduled
    - Previous CPU usage: CPU hogs have lower priority
- ✦ Each thread is placed into a run queue for its priority
  - Head of highest-priority run queue with a ready thread runs next

run queues



# Calculating priority in FreeBSD

- ✦ Thread priority is set by:  
$$\text{pri} = \text{MIN} + [\text{estcpu}/4] + 2 \times \text{nice}$$
  - Values above MAX are set to MAX
  - MIN=160, MAX=223
  - *nice* is set by the user to manually lower thread priority
  - *estcpu* is an estimate of the number of “ready” processes in the CPU when the calculation is made
    - Has a bit of “memory” so it doesn’t change too quickly
    - *estcpu* is updated each clock tick
  - Higher numbers indicate lower priority: threads with lowest priority values are scheduled first
- ✦ Thread priority is set every 40 ms
- ✦ Scheduling is more complex for multiprocessors...

# Scheduling in Linux

---

- ✦ Three classes of processes
  - Conventional
  - Real-time (round-robin)
  - Real-time (FIFO)
- ✦ Queue structure similar to BSD
  - Two sets of queues (0–99 real-time, 100–139 conventional): *active* and *expired*
  - Scheduler runs process in lowest-valued *active* queue
  - Conventional threads placed in *expired* queue when their quantum is up
    - May be scheduled out before quantum expires: put back into *active* queue
  - Real-time threads placed back into *active* queue

# Linux: rescheduling processes

---

- ✦ Re-evaluate priorities when there are no threads on the *active* queue
- ✦ Thread priority is re-evaluated based on previous run time and sleep time
  - Threads that ran more and slept less are penalized with a worse priority
- ✦ Quantum is regenerated based on original (static) priority
  - Higher-priority threads are given longer quanta
  - Process may lose the CPU before quantum is up (if a higher-priority thread becomes available)


# Policy versus mechanism

---

- ✦ Separate what may be done from how it is done
  - Mechanism allows
    - Priorities to be assigned to processes
    - CPU to select processes with high priorities
  - Policy set by what priorities are assigned to processes
- ✦ Scheduling algorithm parameterized
  - Mechanism in the kernel
  - Priorities assigned in the kernel or by users
- ✦ Parameters may be set by user processes
  - Don't allow a user process to take over the system!
  - Allow a user process to voluntarily lower its own priority
  - Allow a user process to assign priority to its threads




# Scheduling user-level threads



- ✦ Kernel picks a process to run next
- ✦ Run-time system (at user level) schedules threads
  - Run each thread for less than process quantum
  - Example: processes get 40ms each, threads get 10ms each
- ✦ Example schedule:  
A1,A2,A3,A1,B1,B3,B2,B3
- ✦ Not possible:  
A1,A2,B1,B2,A3,B3,A2,B1

# Scheduling kernel-level threads



- ✦ Kernel schedules each thread
  - No restrictions on ordering
  - May be more difficult for each process to specify priorities
- ✦ Example schedule:  
A1,A2,A3,A1,  
B1,B3,B2,B3
- ✦ Also possible:  
A1,A2,B1,B2,  
A3,B3,A2,B1